



*PATC*  
*2016/06/06*  
*Maison de la Simulation*

**Understanding and managing  
hardware affinities  
on hierarchical platforms  
With Hardware Locality (hwloc)**

Brice Goglin – TADaaM Team – Inria Bordeaux Sud-Ouest

# Agenda

- Quick example as an Introduction
- Bind your processes
- What's the actual problem?
- Introducing hwloc (Hardware Locality)
- Command-line tools
- C Programming API
- Conclusion

# Materials

- All hwloc tutorials are available at

<http://www.open-mpi.org/projects/hwloc/tutorials>

# 1

**Quick example  
as an Introduction**

# Machines are increasingly complex



# Machines are increasingly complex

- Multiple processors
  - Multicore processors (package = socket)
  - Simultaneous multithreading
  - Shared caches
  - NUMA nodes
  - Multiple GPUs, NICs, ...
- 
- We cannot expect users to understand all this

# Example with MPI

- Our latest cluster at Inria Bordeaux
  - 12-core Xeon E5-2600v3 with NVIDIA K40, etc.
- Nice, let's run some benchmarks!
  - Open MPI 1.8.1
  - Intel MPI benchmarks 3.2

# Example with MPI – Results

- Between cores 0 and 1
  - 540ns, 3500MiB/s
- Between cores 0 and 2
  - 330ns, 4220MiB/s
- Between cores 0 and 12
  - 430ns, 4290MiB/s
- Between cores 0 and 23
  - 590ns, 3410MiB/s



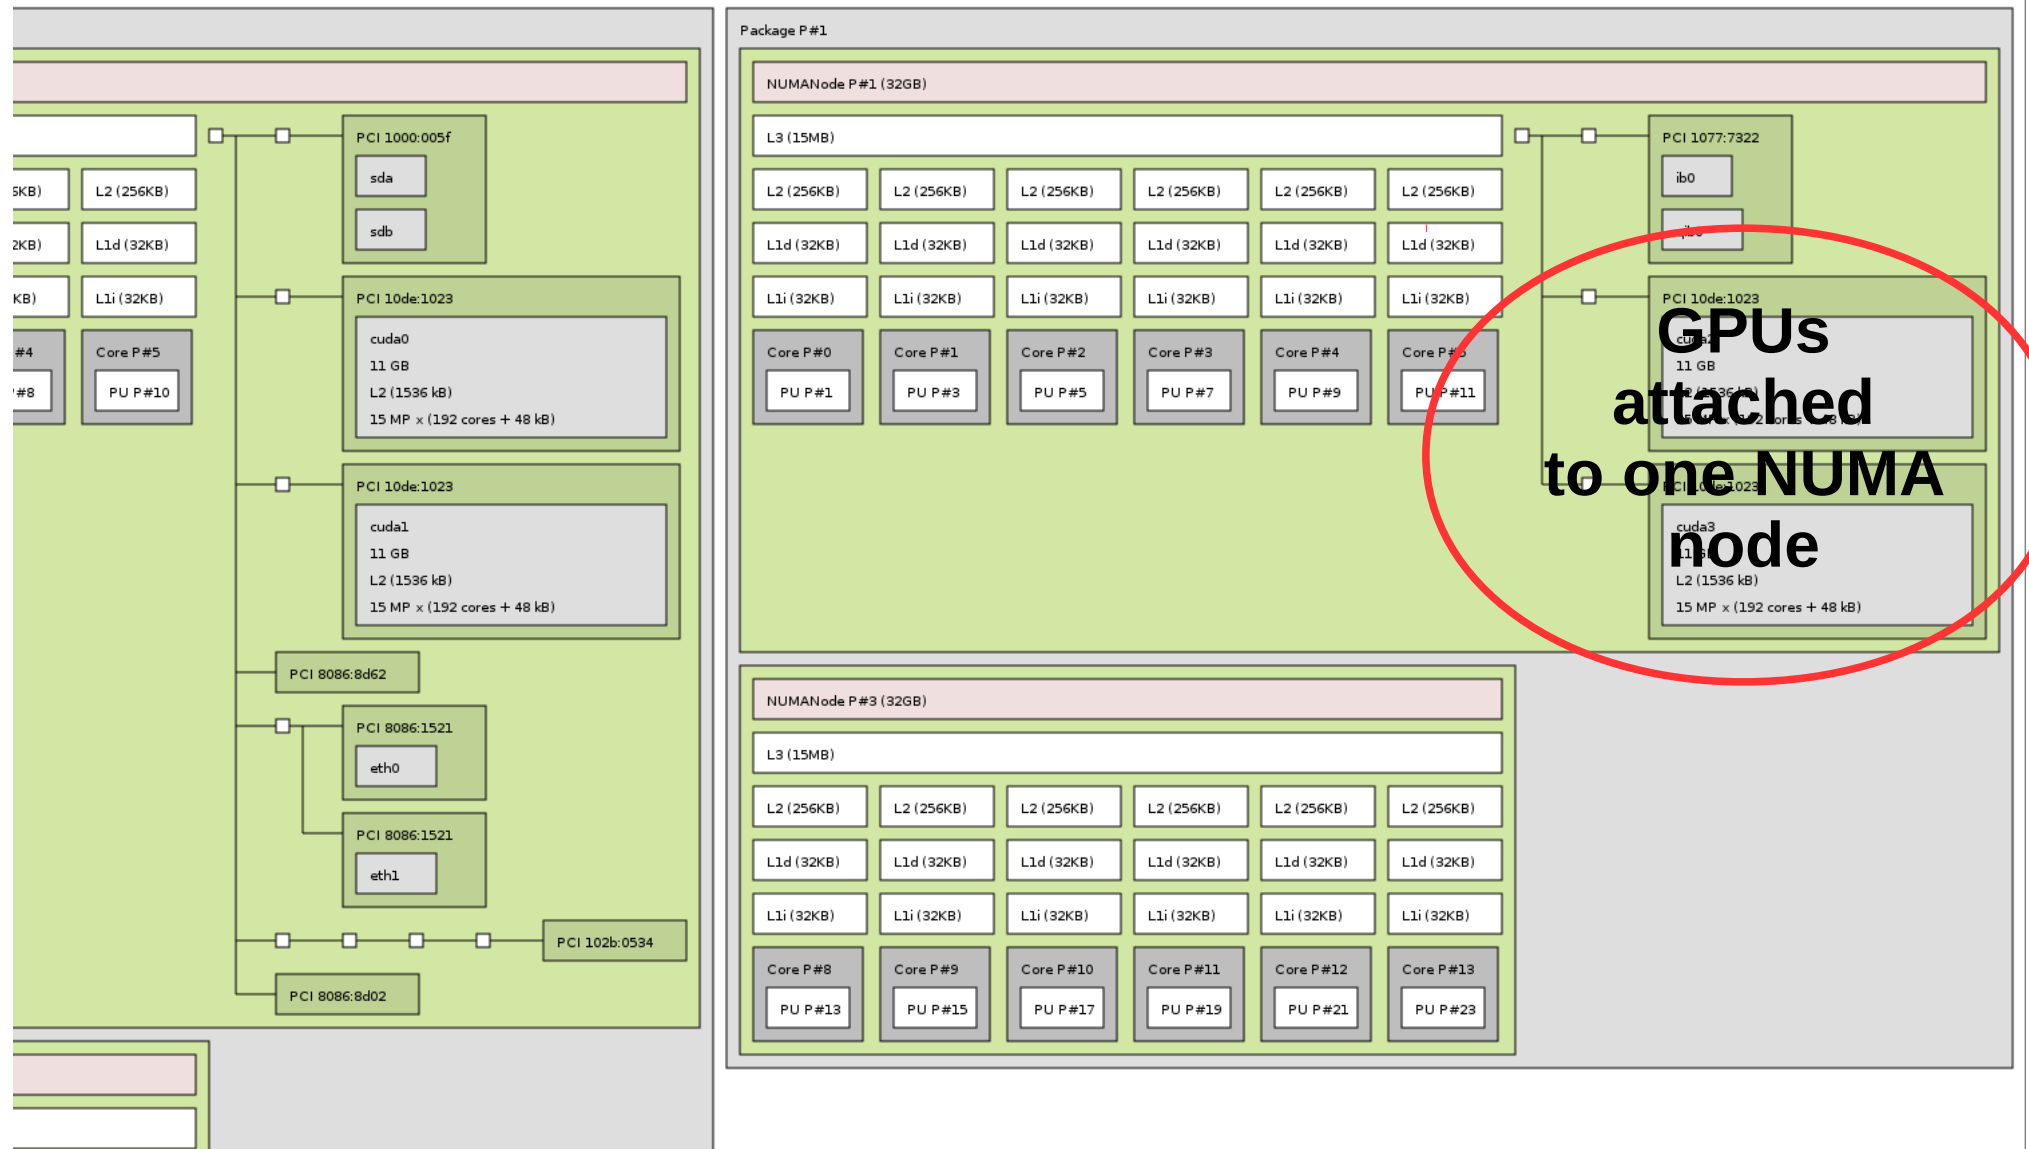
# What is going on?



# First take away messages

- Locality matters to communication performance
  - Machines are really far from flat
    - Similar issues with POWER8, AMD Opterons, Fujitsu Sparc Xlfx
- Cores numbering is crazy
  - Never expect anything sane

# It's actually worse than that



# I/O affinity

- If you use GPUs or high performance networks, you must allocate host memory close to them
  - Otherwise DMA to GPUs slows down by 10-20% here
  - InfiniBand latency increases by 10%
- Need a way to know which cores/memory is close to which I/O device

# 2

## Bind your processes

# Where does locality actually matter?

- MPI communication performance varies with distance
  - Inside or outside nodes
- Shared memory too (threads, OpenMP, etc.)
  - Synchronization
    - Barriers use caches and memory too
  - Concurrent access to shared buffers
    - Producer-consumer, etc
- 15 years ago, locality was mostly an issue for large NUMA SMP machines (SGI, etc.)
  - Today it's everywhere
    - Because multicores and NUMA are everywhere

# What to do about locality in runtimes?


- Place processes/tasks according to their affinities
  - If two tasks communicate/synchronize/share a lot, keep them physically close
    - Main focus of this talk
- Adapt your algorithms to the locality
  - Adapt communication/synchronization implementations to the topology
    - Hierarchical OpenMP barriers
    - Adapt your buffers to (shared) cache size

# Process binding

- Some MPI implementations bind processes by default (Intel MPI, Open MPI 1.8 in some cases)
  - Because it's better for reproducibility
- Some don't
  - Because it may hurt your application
    - Oversubscribing? Dynamic processes?
- Binding doesn't guarantee that your processes are optimally placed
  - It just means your processes won't move
    - No migration, less cache issues, etc.



# To bind or not to bind?




Zeus MHD Blast. 64 Processes/Cores. Mvapich2 1.8. + ICC

# Ok, I need to bind. But where?

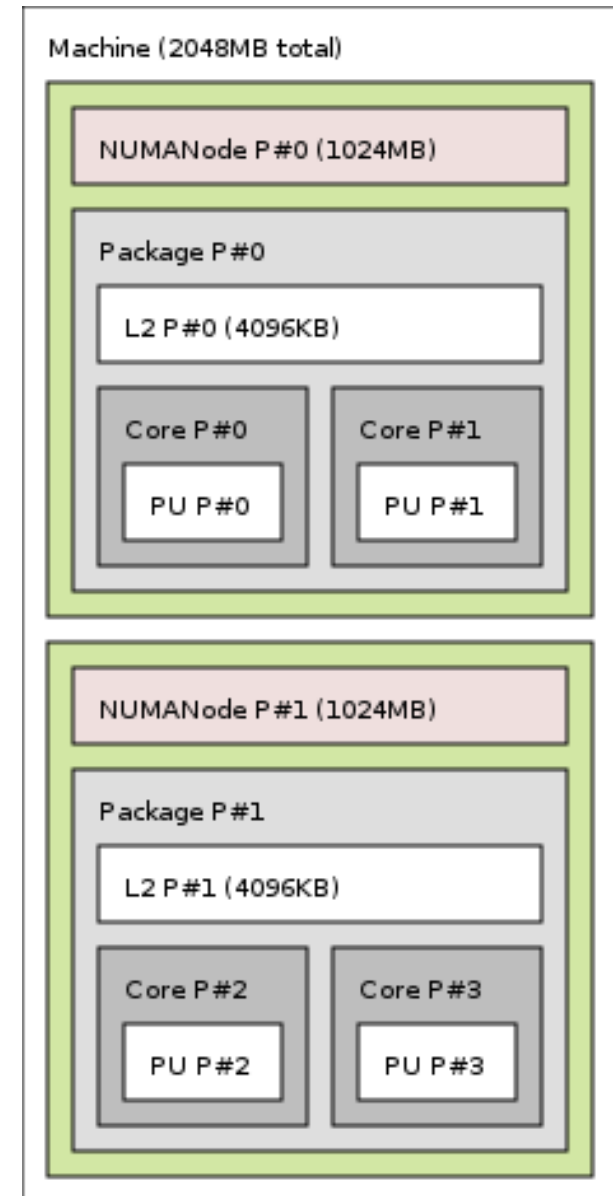
- Default binding strategies?
  - By core first (compact, --map-by core, etc.)
    - One process per core on first node,  
then one process per core on second node, ...
  - By node first (scatter, --map-by node/socket, etc.)
    - One process on first core of each node,  
then one process on second core of each node, ...
- Your application likely prefers one to the other
  - The first one?
    - Because your algorithms often communicate more between  
immediate neighbors
  - Sometimes the other one...

# Binding strategy impact



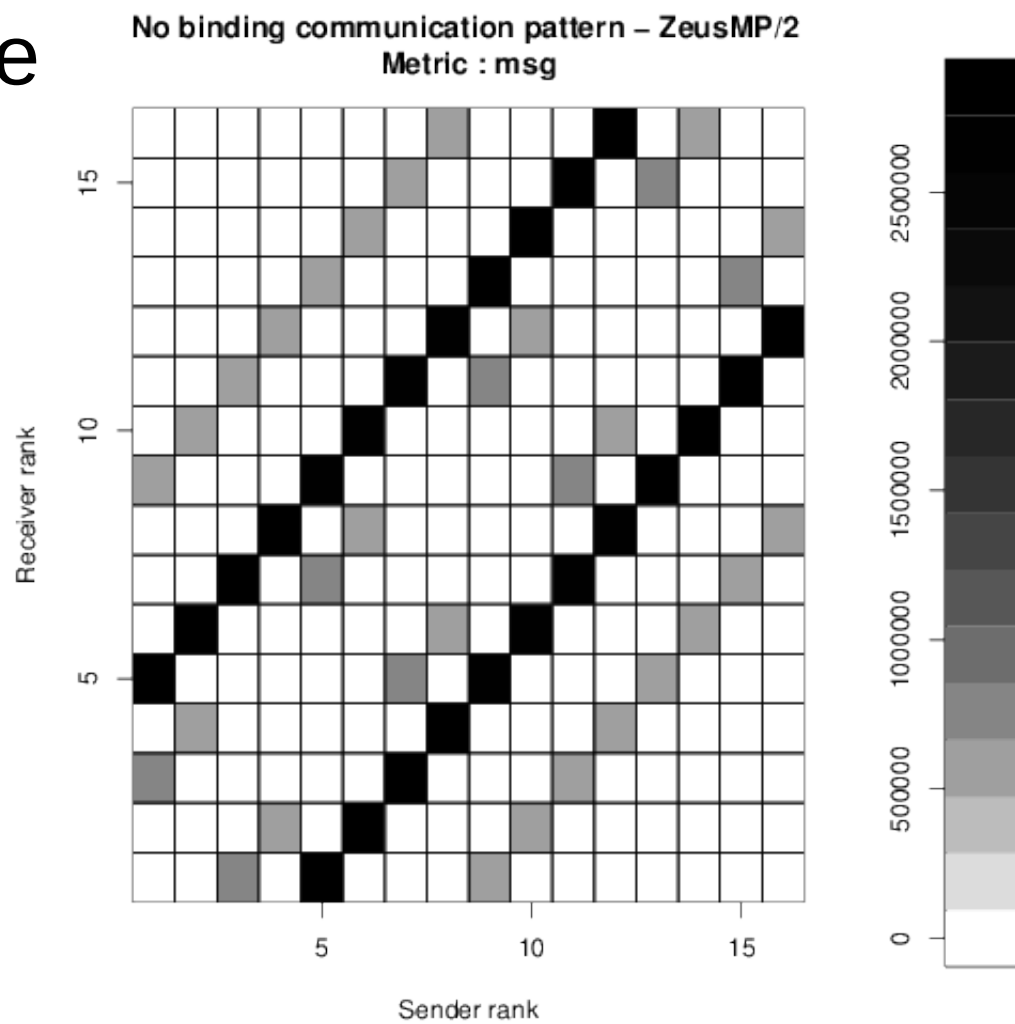
# How do I choose?

- Dilemma
  - Use cores 0 & 1 to share cache and improve synchronization cost?
  - Use cores 0 & 2 to maximize memory bandwidth?
- Depends on the application needs
  - And machine characteristics




# Topology-aware MPI process placement with TreeMatch

- Some tasks communicate a lot with each other
  - The physical distance will slow down some messages
  - Try to keep them close!
- Some don't
  - No constraint on placement




# Reordering tasks to improve locality

No binding communication pattern – ZeusMP/2  
Metric : msg




TreeMatch communication pattern – ZeusMP/2  
Metric : msg



# Reordering with TreeMatch

- At process launch-time
  - mpiexec options
- Dynamically
  - MPI\_Dist\_graph\_create() to swap MPI ranks' roles between application steps
  - Charm++ load-balancer
- The communication volume is unchanged
  - But big volumes move inside nodes
- Faster execution!




# 3

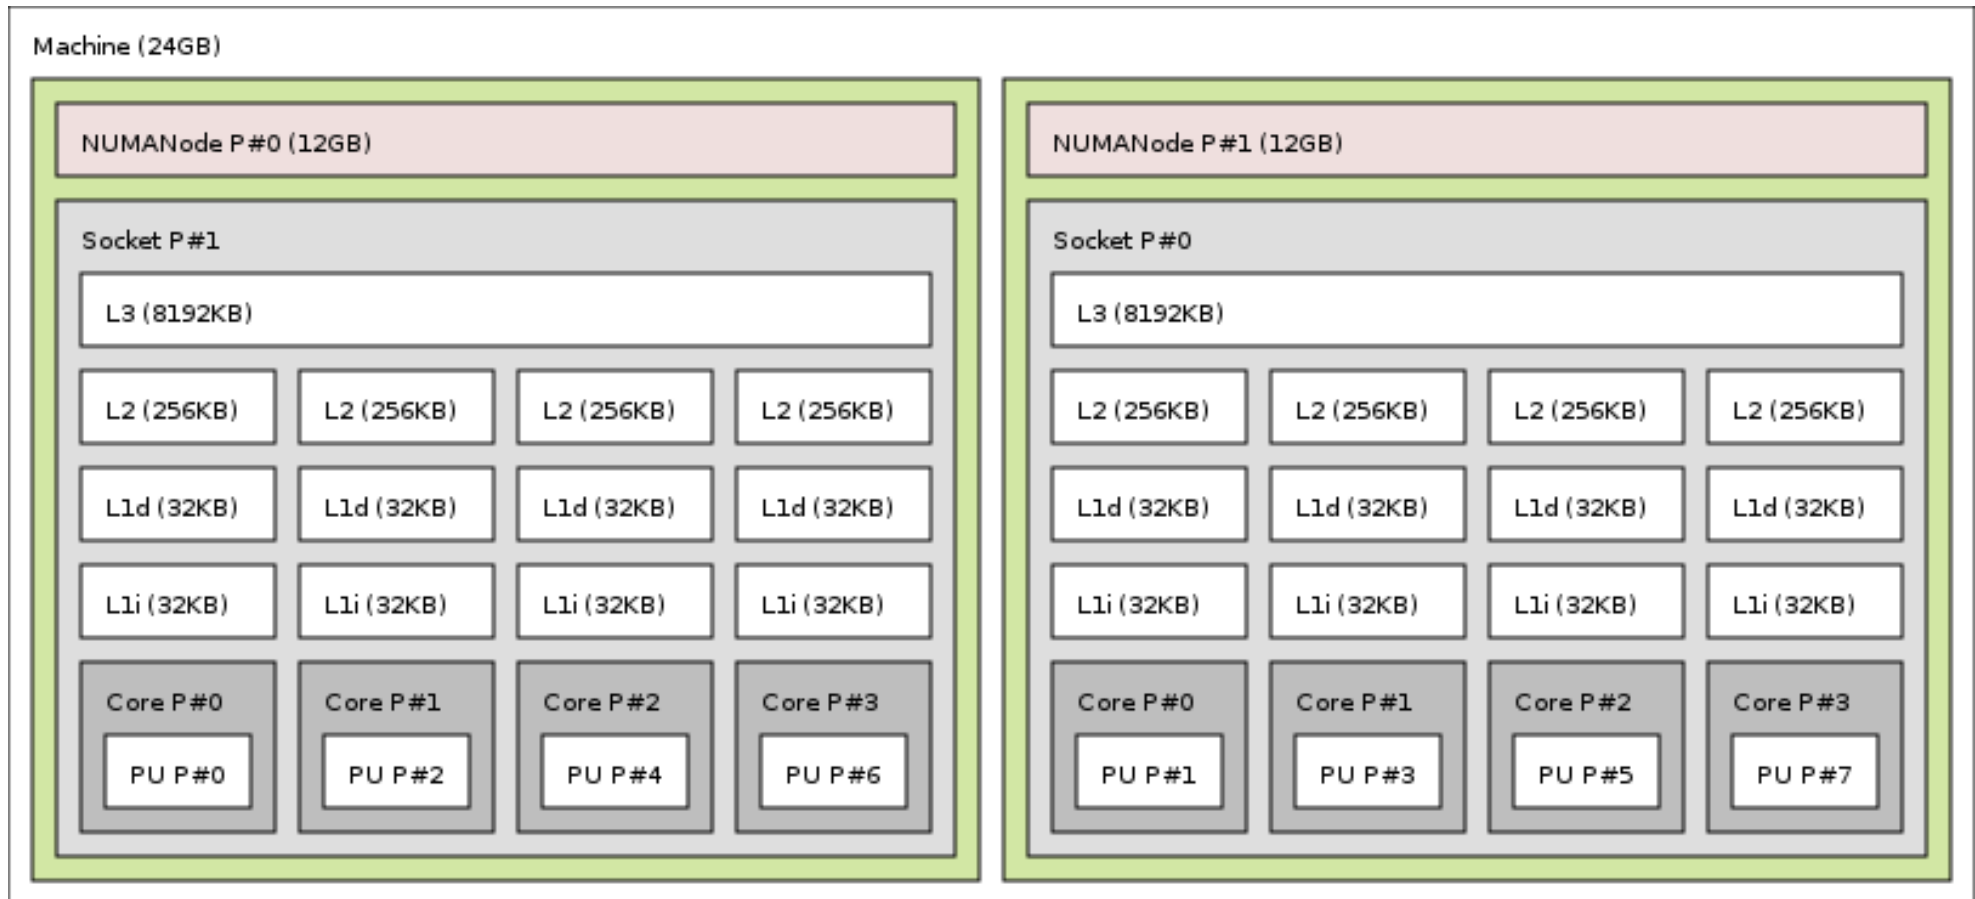
**What's the actual  
problem ?**



# Example of dual Nehalem Xeon machine



# Another example of dual Nehalem Xeon machine



# Processor and core numbers are crazy

- Resource ordering/numbering is unpredictable
  - Can (and does) change with the vendor, BIOS version, etc.
- Some resources may be unavailable
  - Batch schedulers allocates parts of machines
    - Core numbers may be non-consecutive, not start at 0, etc.
- Don't assume anything about these numbers
  - **Otherwise your code won't be portable**

# Vertical ordering of levels (who contains who)



# Vertical ordering isn't reliable either

- Modern processors have 2 NUMA nodes each
  - Xeon E5v3, Opteron 6000, Power8, Sparc64 Xlfx
  - But old platforms have multiple processor packages per NUMA nodes
- Levels of caches and sharing may vary
- Don't assume anything about vertical ordering
  - Or (again) your code won't be portable
  - e.g.: Even the Intel OpenMP binding isn't always correct

# Gathering topology information is difficult

- Lack of generic, uniform interface
  - Operating system specific
    - /proc and /sys on Linux
    - rset, sysctl, lgrp, kstat on other OS
  - Hardware specific
    - x86 CPUID instruction, device-tree, PCI config space, etc.
- Evolving technology
  - AMD Bulldozer introduced dual-core Compute Units
    - It's not two real cores, neither one hyper-threaded core
  - New kinds of hierarchy/resources?
- And some BIOS report buggy information

# Binding is difficult too

- Lack of generic, uniform interface (again)
  - Process/thread binding
    - sched\_affinity() system call changed twice in Linux
  - Memory binding
    - 3 different system-calls on Linux
      - mbind(), migrate\_pages(), move\_pages()
  - Different constraints
    - Bind to single core only? To contiguous set of cores? To random sets of cores?
  - Many different policies

# 4

## Introducing hwloc (Hardware Locality)



# What hwloc is

- Detection of hardware resources
  - Processing units (PU) = logical processors, hardware threads, hyperthreads
    - Things that can run a task
  - Core, packages (sockets), ... (things that contain PUs)
  - Memory nodes, shared caches
  - I/O devices
    - PCI devices and corresponding software handles
- Described as a tree
  - Logical resources identification and organization
    - **Based on locality**

# What hwloc is (2/2)

- API and tools to consult the topology
  - Which cores are near this NUMA memory node ?
  - Give me a single thread in this package
  - Which NUMA memory node is near this GPU ?
  - What shared cache size between these cores ?
- Without caring about hardware strangeness
  - Non portable and crazy numbers, names, ...
- A portable binding API
  - No more Linux sched\_setaffinity() API breakage
  - No more tens of different binding API with different types

# What hwloc is **NOT**

- A placement algorithm
  - hwloc gives hardware information
  - You're the one that knows what your software does/needs
  - You're the one that must match software affinities to hardware localities
    - We give you the hardware information you need
- A performance analysis tool

# hwloc's History

- Ideas from Samuel Thibault's PhD on hierarchical thread scheduling (2003)
- Standalone library to ease MPI process placement (2009)
- Mainly developed by TADaaM@Inria Bordeaux
  - Within the Open MPI consortium
  - Collaboration with many industrial and academic partners
- BSD-3 license
- Many users

# Alternative software with advanced topology knowledge

- numactl/libnuma
  - Only for NUMA + hardware threads
    - No cache, core, package/socket, etc.
- lscpu, lshw, lsusb, ...
  - Specific to some resources
  - Inventory without locality information
- Likwid (performance optimization tool)
  - Now uses hwloc internally
- Intel Compiler (icc)
  - x86 specific, no API

# hwloc Status

- Current stable release : 1.11.3 (April 2016)
- Support for most operating systems and HPC platforms
- Major release every 6 months
  - Backward compatible
- Next major release will be super-major
  - 2.0 expected in 2016H2
  - **Will break the ABI (not fully backward compatible)**
    - Fix bad ideas from the first hwloc API

# hwloc's view of the hardware

- Tree of objects
  - Machines, NUMA memory nodes, packages, caches, cores, threads
    - Logically ordered
  - Grouping similar objects using distances between them
    - Avoids enormous flat topologies
  - Many attributes
    - Memory node size
    - Cache type, size, line size, associativity
    - Physical ordering
    - Miscellaneous info, customizable

# Installing hwloc

- Packages available in Debian, Ubuntu, Redhat, Fedora, CentOS, ArchLinux, NetBSD, etc
- You want the development headers too
  - libhwloc-dev, hwloc-devel, ...



# Manual installation

- Take a recent tarball at <http://www.open-mpi.org/projects/hwloc>
- Dependencies
  - On Linux, numactl/libnuma development headers
  - Cairo headers for lstopo graphics
- `./configure --prefix=$PWD/install`
  - Very few configure options
- Check the summary at the end of configure

# Manual installation

- make
- make install
- Useful environment variables

**export C\_INCLUDE\_PATH=\$C\_INCLUDE\_PATH:<prefix>/include**

**export LD\_LIBRARY\_PATH=\$LD\_LIBRARY\_PATH:<prefix>/lib**

**export PKG\_CONFIG\_PATH=\$PKG\_CONFIG\_PATH:<prefix>/lib/pkgconfig**

**export PATH=\$PATH:<prefix>/bin**

**export MANPATH=\$MANPATH:<prefix>/share/man**

# Using hwloc for this tutorial

- Install hwloc on your preferred cluster
- And install it on your laptop too
  - It will make remote machine consulting easier

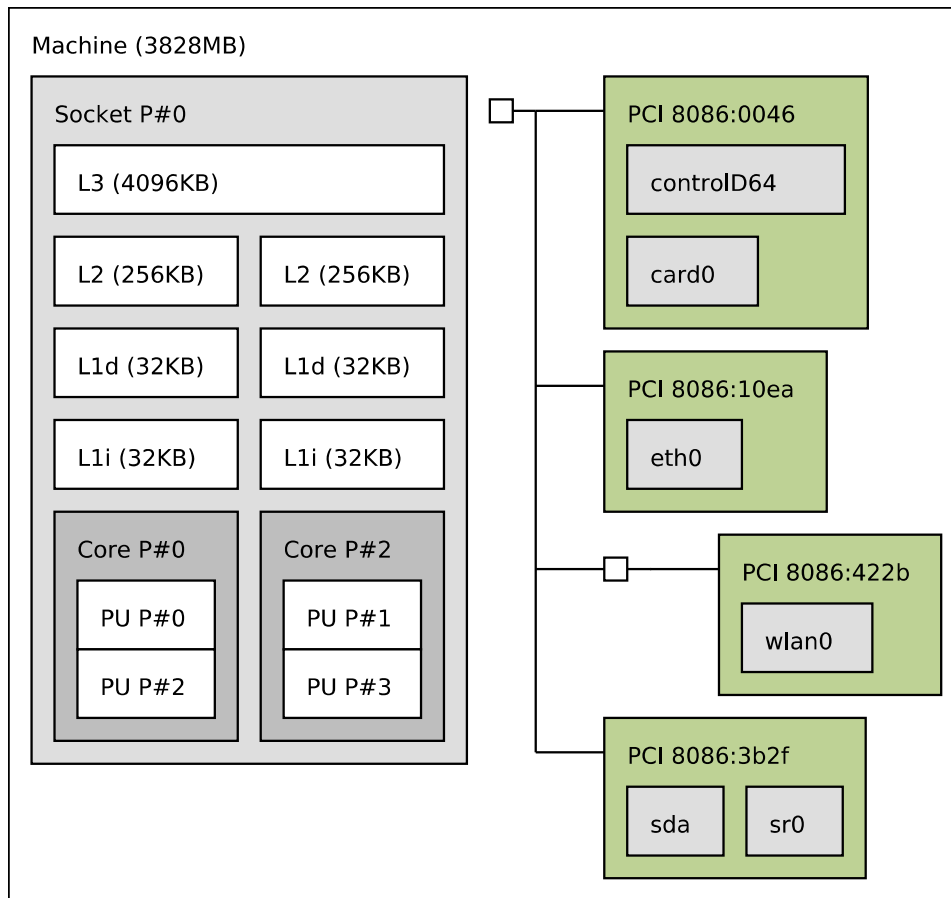
# Using hwloc

- Many hwloc command-line tools
  - lstopo and hwloc-\*
- ... but the actual hwloc power is in the C API
- Perl and Python bindings

# 5

## Command-line Tools

# Istopo (displaying topologies)



Machine (3828MB)

Package L#0 + L3 L#0 (4096KB)

L2 L#0 (256KB) + Core L#0

PU L#0 (P#0)

PU L#1 (P#2)

L2 L#1 (256KB) + Core L#1

PU L#2 (P#1)

PU L#3 (P#3)

HostBridge L#0

PCI 8086:0046

GPU L#0 "controlD64"

PCI 8086:10ea

Net L#2 "eth0"

PCIBridge

PCI 8086:422b

Net L#3 "wlan0"

PCI 8086:3b2f

Block L#4 "sda"

Block L#5 "sr0"

# Istopo

- Many output formats
  - Text, Cairo (PDF, PNG, SVG, PS), Xfig, ncurses
    - Automatically guessed from the file extension
- XML dump/reload
  - Faster, convenient for remote debugging
- Configuration options for nice figures for papers
  - Horizontal/Vertical placement
  - Legend
  - Ignoring some resources
  - Creating fake topologies

# Istopo

```
$ Istopo
```

```
$ Istopo --no-io -
```

```
$ Istopo myfile.png
```

```
$ ssh host Istopo saved.xml
```

```
$ Istopo -i saved.xml
```

```
$ ssh myhost Istopo -.xml | Istopo --if xml -i -
```

```
$ Istopo -i "numa:4 package:2 core:2 pu:2"
```



# hwloc-bind

(binding processes, threads and memory)

- Bind a process to a given set of CPUs  
\$ hwloc-bind package:1 -- mycommand myargs...  
\$ hwloc-bind os=mlx4\_0 -- mympiprogram ...
- Bind an existing process  
\$ hwloc-bind --pid 1234 numa:0
- Bind memory  
\$ hwloc-bind --membind numa:1 --cpubind numa:0 ...
- Find out if a process is already bound  
\$ hwloc-bind --get --pid 1234  
\$ hwloc-ps

# hwloc-calc

## (calculating with objects)

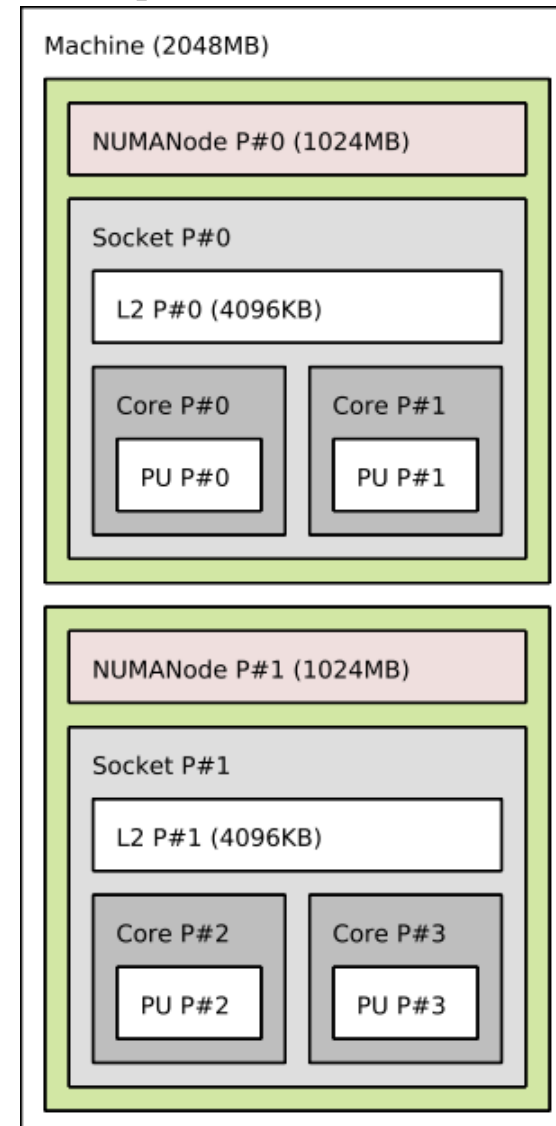
- Convert between ways to designate sets of CPUs, objects... and combine them

```
$ hwloc-calc package:1.core:1 ~pu:even  
0x000000008
```

```
$ hwloc-calc --number-of core numa:0  
2
```

```
$ hwloc-calc --intersect pu package:1  
2,3
```

- The result may be passed to other tools
  - Multiple invocations may be combined
  - I/O devices also supported
- ```
$ hwloc-calc os=eth0
```



# Other tools

- Get some object information
  - hwloc-info
- Generate bitmaps for distributing multiple processes on a topology
  - hwloc-distrib
- Save a Linux node topology info for debugging
  - hwloc-gather-topology
- Manipulating multiple topologies, etc.

# Hands-on Istopo

- Gather the topology of one server
- Display it on another machine
- Hide caches
- Remove the legend
- Restrict the display to a single package
- Export to PDF

# Hands-on hwloc-bind and hwloc-calc

- Bind a process to a core and verify its binding
- Find out how many cores are in the second NUMA node
- Find out the physical numbers of all non-first hyperthreads
- Find out which cores are close to InfiniBand
- Compare the DMA bandwidth from GPU#0 to both NUMA nodes using cudabw

# 6

## C Programming API

# API basics

- A hwloc program looks like this

```
#include <hwloc.h>
```

```
hwloc_topology_t topo;
```

```
hwloc_topology_init(&topo);
```

```
/* ... configure what topology to build ... */
```

```
hwloc_topology_load(topo);
```

```
/* ... play with the topology ... */
```

```
hwloc_topology_destroy(topo);
```

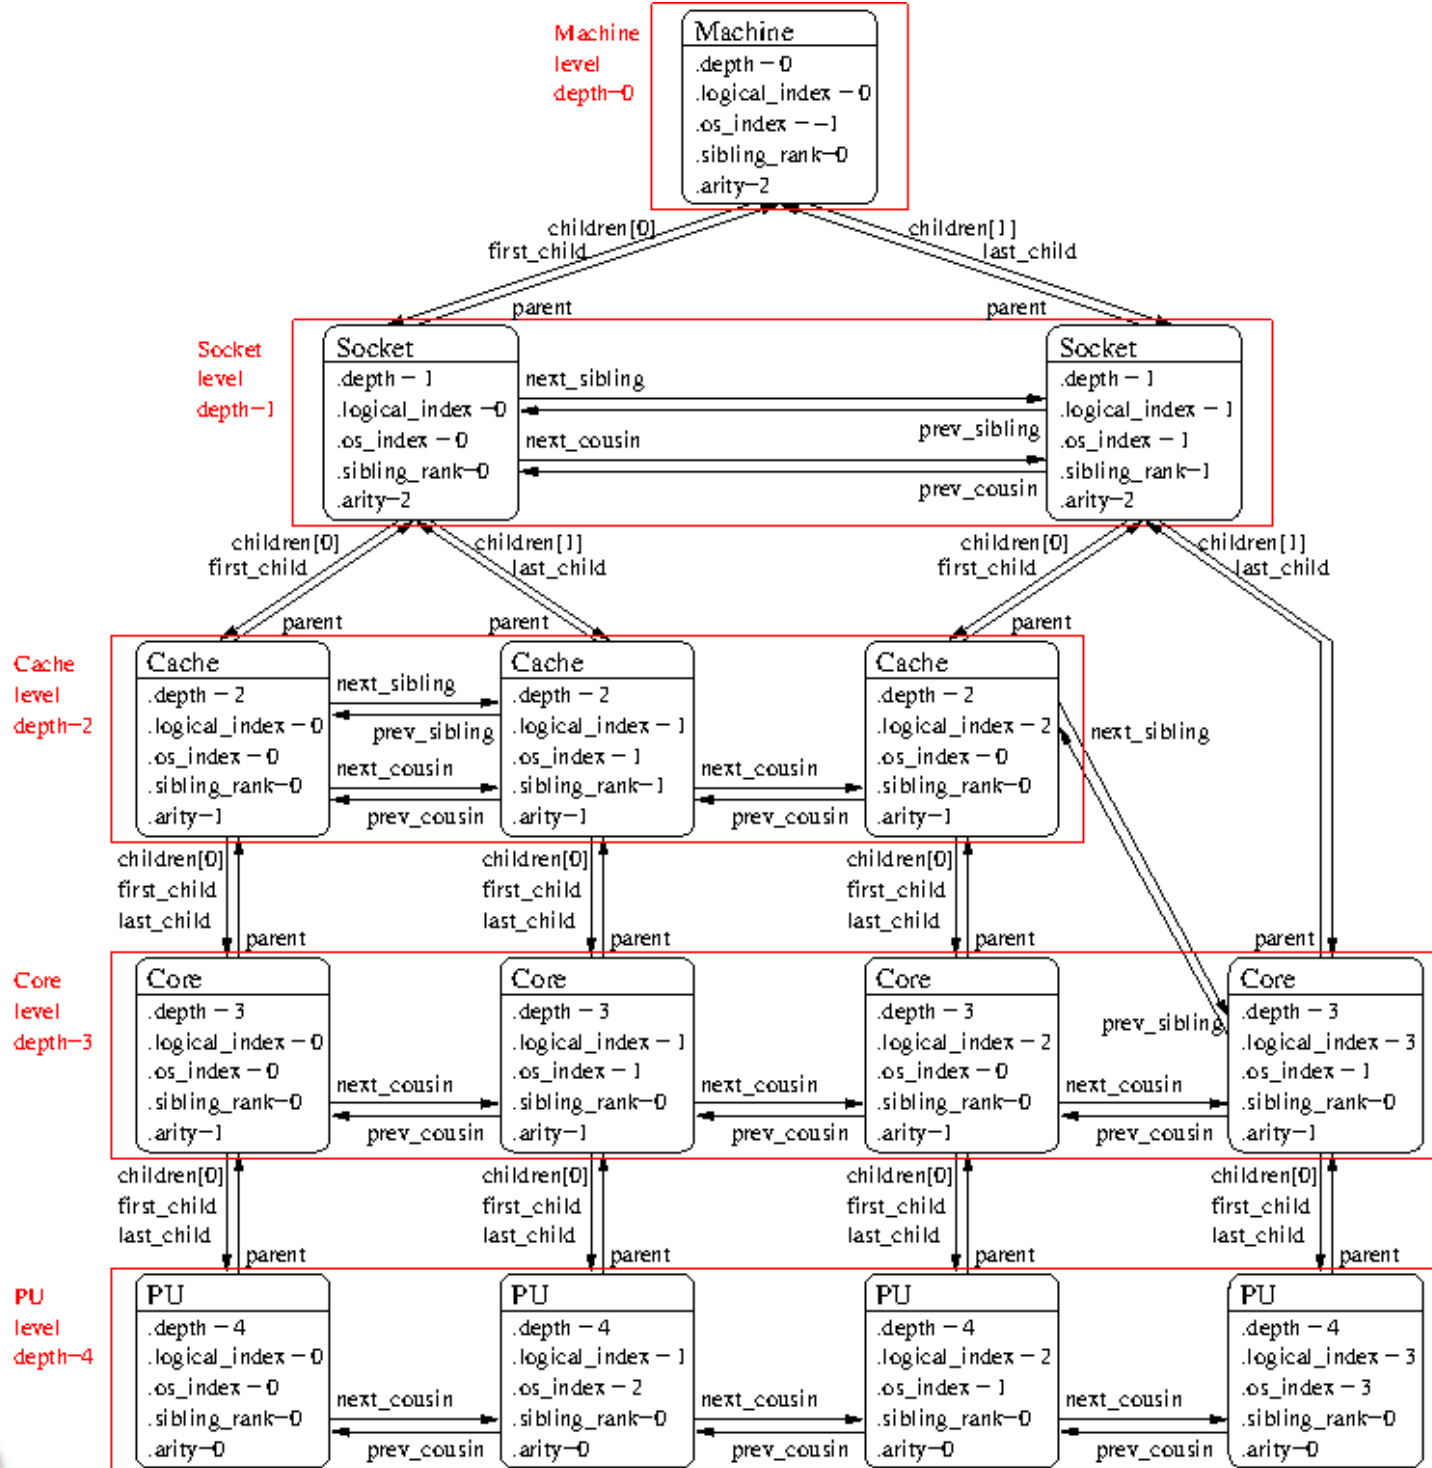
# Major hwloc types

- The topology context : `hwloc_topology_t`
  - You always need one
- The main hwloc object : `hwloc_obj_t`
  - That's where the actual info is
  - The structure isn't opaque
    - It contains many pointers to ease traversal
- Object type : `hwloc_obj_type_t`
  - `HWLOC_OBJ_PU`, `_PACKAGE`, `_CORE`, `_NUMANODE`, ...



# Object information

- Type
- Optional name string
- Indexes (see later)
- cpusets and nodesets (see later)
- Tree pointers (\*cousin, \*sibling, arity, \*child\*, parent)
- Type-specific attribute union
  - obj->attr->cache.size
  - obj->attr->pcidev.linkspeed
- String info pairs



# Browsing as a tree

- The root is `hwloc_get_root_obj(topo)`
- Objects have children
  - `obj->arity` is the number of children
  - The array of children is `obj->children[]`
  - They are also in a list
    - `obj->first_child`, `obj->last_child`
    - `child->prev_sibling`, `child->next_sibling`
    - NULL-terminated
- The parent is `obj->parent` (or NULL)

# Browsing as levels

- The topology is also organized as levels of identical objects
  - Cores, L2d Caches, ...
  - All PUs at the bottom
- Number of levels `hwloc_topology_get_depth(topo)`
- Number of objects on a level  
`hwloc_get_nbobjs_by_type(topo, type)`  
`hwloc_get_nbobjs_by_depth(topo, depth)`
- Convert between depth and type using  
`hwloc_get_type_depth()` or `hwloc_get_depth_type()`

# Browsing as levels

- Find objects by level and index
  - `hwloc_get_obj_by_type(topo, type, index)`
  - There are variants taking a depth instead of a type
    - Note : the depth of my child is not always my depth + 1
      - Think of asymmetric topologies
- Iterate over objects of a level
  - Objects at the same levels are also interconnect by `prev/next_cousin` pointers
    - Don't mix up siblings (children list) and cousins (level)
  - `hwloc_get_next_obj_by_type/depth()`

# Hands-on browsing the topology

Starting from basic.c

- Print the number of cores
- Print the type of the common ancestor of cores 0 and 2
- Print the memory size near core 0
- Iterate over all PUs and print their physical numbers

# Physical or OS indexes

- obj->os\_index
  - The ID given by the OS/hardware
- P#3
  - Default in Istopo graphic mode
  - Istopo -p
- NON PORTABLE
  - Depend on motherboards, BIOS, version, ...
- DON'T USE THEM

# Logical indexes

- obj->logical\_index
  - The index among an entire level
- L#2
  - Default in Istopo except in graphic mode
  - Istopo -l
- Always represent proximity (depth-first walk)
- PORTABLE
  - Does not depend on OS/BIOS/weather
- That's what you want to use



# But I still need OS indexes when binding ?!

- NO !
- Just use hwloc for binding, you won't need physical/OS indexes ever again
- If you want to bind the execution to a core
  - `hwloc_set_cpubind(core->cpuset)`
    - Other API functions for binding entire processes, single thread, memory, for allocating bound memory, etc.

# Bitmap, CPU sets, Node sets

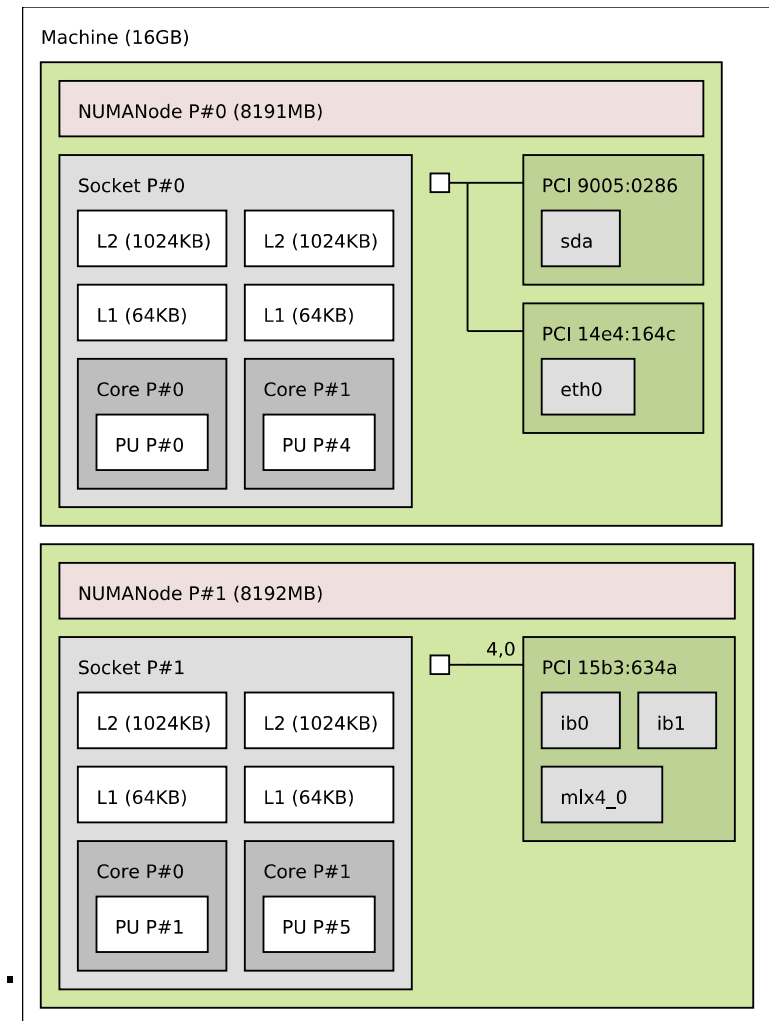
- Generic mask of bits : `hwloc_bitmap_t`
  - Possibly infinite
  - Opaque, used to describe object contents
    - Which PU are inside this object (`obj->cpuset`)
    - Which NUMA nodes are close to this object (`obj->nodeset`)
  - Can be combined to bind to multiple cores, etc.
    - and, or, xor, not, ...

# Hands-on bitmaps and binding

- Bind a process to 1<sup>st</sup> core
- Rebind the same process to cores 2 and 4
- Print its binding
- Print where it's actually running
  - Repeat
- Rebind to avoid migrating between cores
  - `hwloc_bitmap_singlify()`

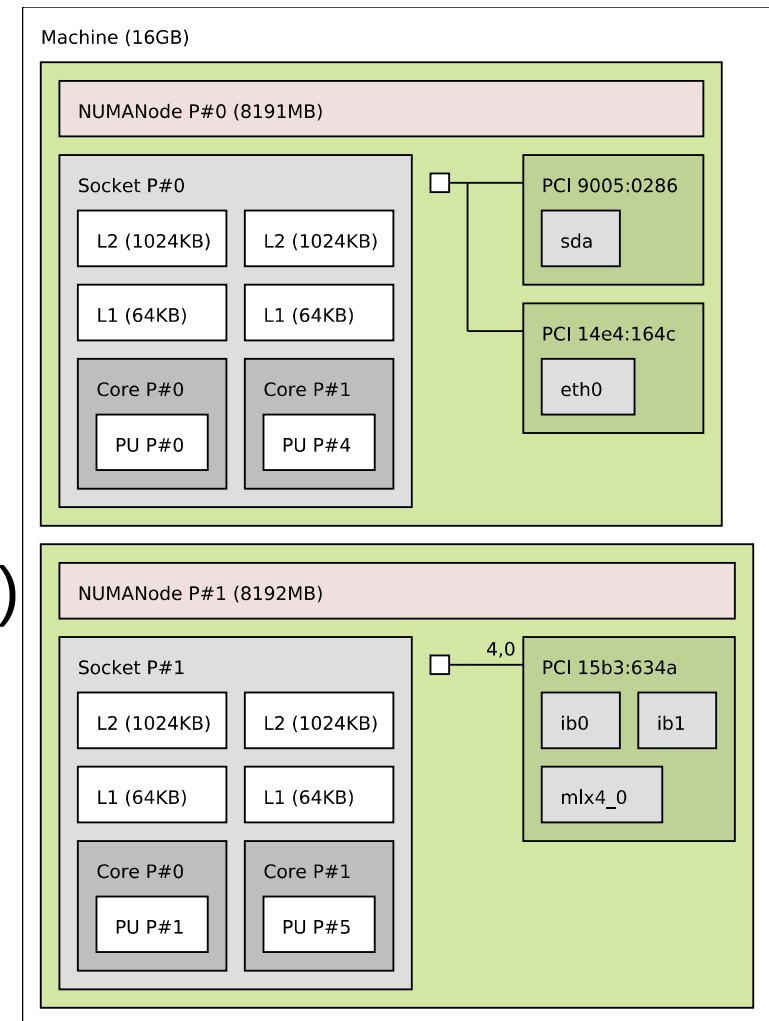
# I/O devices

- Binding tasks near the devices they use improves their data transfer time
  - GPUs, high-performance NICs, InfiniBand, ...
- You cannot bind tasks or memory on these devices
  - But these devices may have interesting attributes
    - Device type, GPU capabilities, embedded memory, link speed, ...



# I/O objects

- Some I/O trees are attached to the object they are close to
- PCI device objects
  - Optional I/O bridge objects
- How to match your software handle with a PCI device ?
  - OS/Software devices (when known)
    - sda, eth0, ib0, mlx4\_0
- Disabled by default
  - Except in Istopo



# Hands-on I/O

- Load cuda modules, etc.

Starting from cuda.c

- Find the NUMA node near each CUDA device

Starting from ib.c

- Find the NUMA node near each IB device

# Extended attributes

- obj->userdata pointer
  - Your application may store whatever it needs there
  - hwloc won't look at it, it doesn't know what's it contains
- (name,value) info attributes
  - Basic string annotations, hwloc adds some
    - Hostname, Kernel release, CPU model, PCI vendor, ...
    - See lstopo -v for (many) examples
  - You may add your own

# Configuring the topology

- Between `hwloc_topology_init()` and `load()`
  - `hwloc_topology_set_xml()`, `set_synthetic()`
  - `hwloc_topology_set_flags()`, `set_pid()`
  - `hwloc_topology_ignore_type()`
- After `hwloc_topology_load()`
  - `hwloc_topology_restrict()`
  - `hwloc_topology_insert_misc_object...`



# Helpers

- hwloc/helper.h contains a lot of helper functions
  - Iterators on levels, children, restricted levels
  - Finding caches
  - Converting between cpusets and nodesets
  - Finding I/O objects
  - And much more
- Use them to avoid rewriting basic functions
- Use them to understand how things work and write what you need

# 8

## Conclusion

# More information

- The documentation
  - <http://www.open-mpi.org/projects/hwloc/doc/>
  - Related pages
    - <http://www.open-mpi.org/projects/hwloc/doc/v1.11.3/pages.php>
  - FAQ
    - <http://www.open-mpi.org/projects/hwloc/doc/v1.11.3/a00030.php>
- README and HACKING in the source
- hwloc-users@open-mpi.org for questions
- hwloc-devel@open-mpi.org for contributing
- hwloc-announce@open-mpi.org for new releases
- <https://github.com/open-mpi/hwloc/issues> for reporting bugs

Thanks!

Questions?

<http://www.open-mpi.org/projects/hwloc>



Brice.Goglin@inria.fr